

Assignment 4 (M1)

```

✓ [94] # Loading an audio file
0s      import warnings
        warnings.filterwarnings('ignore')

        import librosa

```

```

✓ [3]  from google.colab import files
13m      uploaded = files.upload()

```

No files selected. Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving genres.zip to genres.zip

```

✓ [73] !unzip genres.zip
18s

```

Archive: genres.zip
replace genres/pop/only_the_young_chorus.wav? [y]es, [n]o, [A]ll, [N]one, [r]ename: n
replace genres/rock/dont_stop_me_now_chorus.wav? [y]es, [n]o, [A]ll, [N]one, [r]ename: n

```


✓ [155] #####
1s      # First wav music file
        audio_path = 'genres/pop/only_the_young_chorus.wav'
        x , sr = librosa.load(audio_path)

```

```

✓ [98] # Playing Audio
6s      # Using IPython.display.Audio, to play the audio
        import IPython.display as ipd
        ipd.Audio(audio_path)

```




```

✓ [99] # Visualizing Audio
1s      ## 1. Waveform ##
        %matplotlib inline
        import sklearn
        import matplotlib.pyplot as plt
        import librosa.display

        plt.figure(figsize=(20, 5))
        librosa.display.waveplot(x, sr=sr)

```

<matplotlib.collections.PolyCollection at 0x7f8970e95490>




```

✓ [100] ## 2. Spectrogram ##
2s
x = librosa.stft(x)
Xdb = librosa.amplitude_to_db(abs(X))
plt.figure(figsize=(20, 5))
librosa.display.specshow(Xdb, sr=sr, x_axis='time', y_axis='hz')
plt.colorbar()

```

<matplotlib.colorbar.Colorbar at 0x7f897036cf10>




```

✓ [101] ## 3. Log Frequency axis ##
1s
librosa.display.specshow(Xdb, sr=sr, x_axis='time', y_axis='log')
plt.colorbar()

```


<matplotlib.colorbar.Colorbar at 0x7f89707b95d0>



```

✓ [136] # Feature Extraction
0s
## 1. Zero Crossing Rate ##
# Zooming in
# Here we will zoom or print spectrum for 100 array columns only.
n0 = 9000
n1 = 9100
plt.figure(figsize=(20, 5))
plt.plot(x[n0:n1])
plt.grid()

```



```
✓ [141] zero_crossings = librosa.zero_crossings(x[n0:n1], pad=False)
0s zero_crossings.shape
```

```
(100,)
```

```
✓ [142] # We can also calculate zero crossings using a given code:
0s print(sum(zero_crossings))
```

```
31
```

```
✓ [145] ## 2. Spectral Centroid ##
0s #spectral centroid -- centre of mass -- weighted mean of the frequencies present in the sound
spectral_centroids = librosa.feature.spectral_centroid(x, sr=sr)[0]
spectral_centroids.shape
```


```
(1034,)
```

```
✓ [146] # Computing the time variable for visualization
1s plt.figure(figsize=(20,5))
frames = range(len(spectral_centroids))
t = librosa.frames_to_time(frames)

# Normalising the spectral centroid for visualisation
def normalize(x, axis=0):
    return sklearn.preprocessing.minmax_scale(x, axis=axis)


# Plotting the Spectral Centroid along the waveform
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_centroids), color='r')
```

```
[<matplotlib.lines.Line2D at 0x7f896fa2d150>]
```



```
✓ [147] ## 3. Spectral Bandwidth ##
2s plt.figure(figsize=(20,5))
spectral_bandwidth_2 = librosa.feature.spectral_bandwidth(x+0.01, sr=sr)[0]
spectral_bandwidth_3 = librosa.feature.spectral_bandwidth(x+0.01, sr=sr, p=3)[0]
spectral_bandwidth_4 = librosa.feature.spectral_bandwidth(x+0.01, sr=sr, p=4)[0]
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_bandwidth_2), color='r')
plt.plot(t, normalize(spectral_bandwidth_3), color='g')
plt.plot(t, normalize(spectral_bandwidth_4), color='y')
plt.legend(('p = 2', 'p = 3', 'p = 4'))
```


```
<matplotlib.legend.Legend at 0x7f896fa07f10>
```




```

✓ [149] ## 4. Spectral Rolloff ##
1s plt.figure(figsize=(20,5))
spectral_rolloff = librosa.feature.spectral_rolloff(x+0.01, sr=sr)[0]
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_rolloff), color='r')
plt.grid()

```




```

✓ [150] ## 5. Mel-Frequency Cepstral Coefficients ##
1s plt.figure(figsize=(20,5))
mfccs = librosa.feature.mfcc(x, sr=sr)
print(mfccs.shape)

#Displaying the MFCCs:
librosa.display.specshow(mfccs, sr=sr, x_axis='time')

(20, 1034)
<matplotlib.collections.QuadMesh at 0x7f897044e750>

```



```

✓ [151] # MFCC Feature Scaling
0s # Let's scale the MFCCs such that each coefficient dimension has zero mean and unit variance:
mfccs = sklearn.preprocessing.scale(mfccs, axis=1)
print(mfccs.mean(axis=1))
print(mfccs.var(axis=1))

```

```


[ 7.8396827e-09  0.0000000e+00 -1.4757050e-08  7.3785249e-09
 -7.3785249e-09  7.3785249e-09 -7.3785249e-09  1.1067787e-08
 -7.3785249e-09  0.0000000e+00  0.0000000e+00 -5.5338933e-09
 -7.3785249e-09  7.3785249e-09 -7.3785249e-09  0.0000000e+00
  0.0000000e+00  7.3785249e-09  0.0000000e+00  0.0000000e+00]

[1.  1.  1.  1.  0.9999999 1.  1.
 1.  1.0000001 1.  1.  1.  1.
0.9999999 1.  1.  1.  1.0000001 1.]

```


```
✓ [153] plt.figure(figsize=(20,8))
0s librosa.display.specshow(mfccs, sr=sr, x_axis='time')
```

<matplotlib.collections.QuadMesh at 0x7f896fb63210>




```
✓ [180] ## 6. Chroma Frequencies ##
1s hop_length = 512
chromagram = librosa.feature.chroma_stft(x, sr=sr, hop_length=hop_length)
plt.figure(figsize=(20, 5))
librosa.display.specshow(chromagram, x_axis='time', y_axis='chroma', hop_length=hop_length, cmap='coolwarm')
```

<matplotlib.collections.QuadMesh at 0x7f896f2840d0>




```
✓ [157] ## 7. Root-mean-square energy (RMSE) ##
1s plt.figure(figsize=(20,5))
hop_length = 256
frame_length = 512
rmse = librosa.feature.rms(x, frame_length=frame_length, hop_length=hop_length, center=True)[0]
frames = range(len(rmse))
t = librosa.frames_to_time(frames, sr=sr, hop_length=hop_length)
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t[:len(rmse)], rmse/rmse.max(), color='green') # normalized for visualization
plt.legend(["RMSE"], loc=0, frameon=True)
```

<matplotlib.legend.Legend at 0x7f896ff5a8d0>




```
✓ [158] #####  
1s # Second wav music file  
audio_path = 'genres/rock/dont_stop_me_now_chorus.wav'  
x, sr = librosa.load(audio_path)
```

```
✓ [159] # Playing Audio  
6s # Using IPython.display.Audio, to play the audio  
import IPython.display as ipd  
ipd.Audio(audio_path)
```




```
✓ [160] # Visualizing Audio  
1s ## 1. Waveform ##  
%matplotlib inline  
import sklearn  
import matplotlib.pyplot as plt  
import librosa.display  
  
plt.figure(figsize=(20, 5))  
librosa.display.waveplot(x, sr=sr)
```

<matplotlib.collections.PolyCollection at 0x7f8971258650>

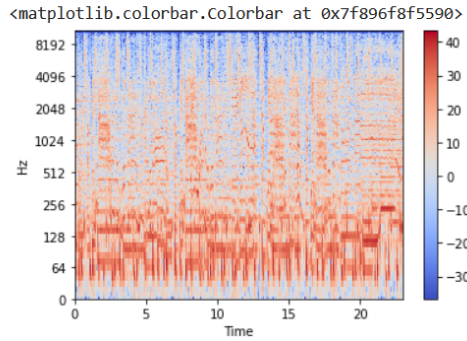


```
✓ [161] ## 2. Spectrogram ##  
3s x = librosa.stft(x)  
Xdb = librosa.amplitude_to_db(abs(X))  
plt.figure(figsize=(20, 5))  
librosa.display.specshow(Xdb, sr=sr, x_axis='time', y_axis='hz')  
plt.colorbar()
```


<matplotlib.colorbar.Colorbar at 0x7f896f8c8310>



```
✓ [162] ## 3. Log Frequency axis ##
2s librosa.display.specshow(Xdb, sr=sr, x_axis='time', y_axis='log')
plt.colorbar()
```



```
✓ [163] # Feature Extraction
0s ## 1. Zero Crossing Rate ##
# Zooming in
# Here we will zoom or print spectrum for 100 array columns only.
n0 = 9000
n1 = 9100
plt.figure(figsize=(20, 5))
plt.plot(x[n0:n1])
plt.grid()
```



```
✓ [164] zero_crossings = librosa.zero_crossings(x[n0:n1], pad=False)
0s zero_crossings.shape

(100,)
```

```
✓ [165] # We can also calculate zero crossings using a given code:
0s print(sum(zero_crossings))
```

8

```
✓ [166] ## 2. Spectral Centroid ##
0s #spectral centroid -- centre of mass -- weighted mean of the frequencies present in the sound
spectral_centroids = librosa.feature.spectral_centroid(x, sr=sr)[0]
spectral_centroids.shape
```

(991,)


```


✓ [167] # Computing the time variable for visualization
15 plt.figure(figsize=(20,5))
frames = range(len(spectral_centroids))
t = librosa.frames_to_time(frames)

# Normalising the spectral centroid for visualisation
def normalize(x, axis=0):
    return sklearn.preprocessing.minmax_scale(x, axis=axis)

# Plotting the Spectral Centroid along the waveform
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_centroids), color='r')

```

[<matplotlib.lines.Line2D at 0x7f896f834bd0>]




```

✓ [168] ## 3. Spectral Bandwidth ##
15 plt.figure(figsize=(20,5))
spectral_bandwidth_2 = librosa.feature.spectral_bandwidth(x+0.01, sr=sr)[0]
spectral_bandwidth_3 = librosa.feature.spectral_bandwidth(x+0.01, sr=sr, p=3)[0]
spectral_bandwidth_4 = librosa.feature.spectral_bandwidth(x+0.01, sr=sr, p=4)[0]
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_bandwidth_2), color='r')
plt.plot(t, normalize(spectral_bandwidth_3), color='g')
plt.plot(t, normalize(spectral_bandwidth_4), color='y')
plt.legend(('p = 2', 'p = 3', 'p = 4'))

```


<matplotlib.legend.Legend at 0x7f896f7a9ed0>



```

✓ [169] ## 4. Spectral Rolloff ##
15 plt.figure(figsize=(20,5))
spectral_rolloff = librosa.feature.spectral_rolloff(x+0.01, sr=sr)[0]
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_rolloff), color='r')
plt.grid()

```




```

✓ [170] ## 5. Mel-Frequency Cepstral Coefficients ##
Ds
plt.figure(figsize=(20,5))
mfccs = librosa.feature.mfcc(x, sr=sr)
print(mfccs.shape)


#Displaying the MFCCs:
librosa.display.specshow(mfccs, sr=sr, x_axis='time')

```

```

(20, 991)
<matplotlib.collections.QuadMesh at 0x7f896f6a9f10>

```



```

✓ [171] # MFCC Feature Scaling
Ds
# Let's scale the MFCCs such that each coefficient dimension has zero mean and unit variance:
mfccs = sklearn.preprocessing.scale(mfccs, axis=1)
print(mfccs.mean(axis=1))
print(mfccs.var(axis=1))

```

```

[ 0.0000000e+00  0.0000000e+00  0.0000000e+00 -1.9246706e-09
 7.6986826e-09  0.0000000e+00  0.0000000e+00 -7.6986826e-09
 7.6986826e-09  0.0000000e+00  7.6986826e-09  3.8493413e-09
 0.0000000e+00  0.0000000e+00  0.0000000e+00  7.6986826e-09
 5.7740119e-09  9.6233537e-09  1.1548024e-08  2.8870060e-09]
[1.  1.  1.0000001  0.9999999  1.0000001  1.0000001
 1.0000001  0.99999994  1.0000001  0.9999999  1.  1.0000001
 1.0000001  1.  0.9999999  1.  1.  0.9999998
 1.  1.0000001 ]

```

```


✓ [172] plt.figure(figsize=(20,8))
Ds
librosa.display.specshow(mfccs, sr=sr, x_axis='time')

```

```

<matplotlib.collections.QuadMesh at 0x7f896f678910>

```




```

✓ [175] ## 6. Chroma Frequencies ##
0s
hop_length = 512
chromagram = librosa.feature.chroma_stft(x, sr=sr, hop_length=hop_length)
plt.figure(figsize=(20, 5))
librosa.display.specshow(chromagram, x_axis='time', y_axis='chroma', hop_length=hop_length, cmap='coolwarm')

```

<matplotlib.collections.QuadMesh at 0x7f896f4b0990>




```

✓ [176] ## 7. Root-mean-square energy (RMSE) ##
1s
plt.figure(figsize=(20,5))
hop_length = 256
frame_length = 512
rmse = librosa.feature.rms(x, frame_length=frame_length, hop_length=hop_length, center=True)[0]
frames = range(len(rmse))
t = librosa.frames_to_time(frames, sr=sr, hop_length=hop_length)
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t[:len(rmse)], rmse/rmse.max(), color='green') # normalized for visualization
plt.legend(["RMSE"], loc=0, frameon=True)

```

<matplotlib.legend.Legend at 0x7f896f40fd10>



```

✓ [184] #####
0s
# Creating the CSV Dataset
import os
import csv
import numpy as np

header = 'filename chroma_stft rmse spectral_centroid spectral_bandwidth rolloff zero_crossing_rate'
for i in range(1, 21):
    header += f' mfcc{i}'
header += ' label'
header = header.split()

```

```
✓ [186] file = open('data.csv', 'w', newline='')
3s
with file:
    writer = csv.writer(file)
    writer.writerow(header)
    genres = 'pop rock'.split()

    for g in genres:
        for filename in os.listdir(f'./genres/{g}'):
            songname = f'./genres/{g}/{filename}'
            y, sr = librosa.load(songname, mono=True, duration=30)
            chroma_stft = librosa.feature.chroma_stft(y=y, sr=sr)
            rmse = librosa.feature.rms(y=y)
            spec_cent = librosa.feature.spectral_centroid(y=y, sr=sr)
            spec_bw = librosa.feature.spectral_bandwidth(y=y, sr=sr)
            rolloff = librosa.feature.spectral_rolloff(y=y, sr=sr)
            zcr = librosa.feature.zero_crossing_rate(y)
            mfcc = librosa.feature.mfcc(y=y, sr=sr)
            to_append = f'{filename} {np.mean(chroma_stft)} {np.mean(rmse)} {np.mean(spec_cent)} {np.mean(spec_bw)} {np.mean(rolloff)} {np.mean(zcr'
            for e in mfcc:
                to_append += f' {np.mean(e)}'
            to_append += f' {g}'
            file = open('data.csv', 'a', newline='')
            with file:
                writer = csv.writer(file)
                writer.writerow(to_append.split())
```